




Figure 1: Model error versus policy shift on Hopper. The 50k-step dynamics model is used to evaluate prediction error, and policy shift is measured as the KL divergence between the 50k-step policy and historical policies.




Figure 2: β ablation (Hopper).




Figure 3: β ablation (Humanoid).




Figure 4: λ ablation (Humanoid).




Figure 5: Empirical estimate of $\delta_{\hat{T}_\phi}$ versus training size.




Figure 6: Empirical estimate of $L_s^{(k)}(\hat{T}_\phi)$ on Hopper. We estimate the $k=10$ consistency coefficient by discretizing the state space into bins, constructing the policy-induced closed-loop transition kernel using the learned dynamics model, computing its k -step kernel P^k , and evaluating $\delta(P^k) = \frac{1}{2} \max_{i,j} \|P^k(i, \cdot) - P^k(j, \cdot)\|_1$. Lower values indicate better finite-horizon multi-step consistency.




Figure 7: Comparison between PMER and MaPER-MBPO on the four benchmark environments used in MaPER.




Figure 8: Comparison of PMER, the uniform baseline, and a random-weight proxy on Humanoid task.




Figure 9: Meta-World results with a Dreamer-v3 backbone.