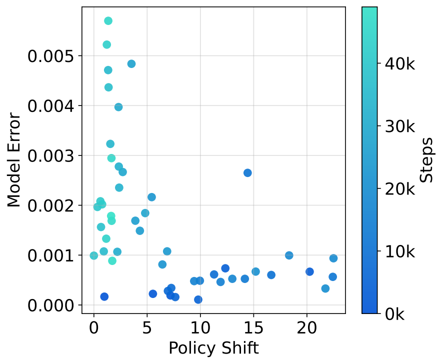

Figure 1: Model error versus policy shift on Hopper. The model checkpoint at 50k training steps is used as the reference for measuring distances to historical policies.

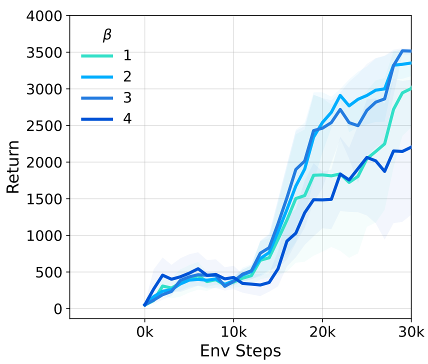

Figure 2: β ablation (Hopper).

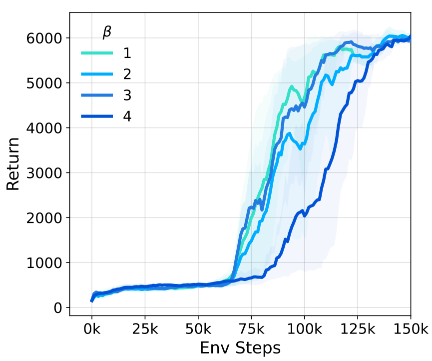

Figure 3: β ablation (Humanoid).

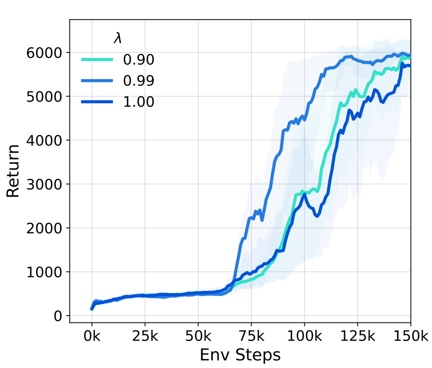

Figure 4: λ ablation (Humanoid).

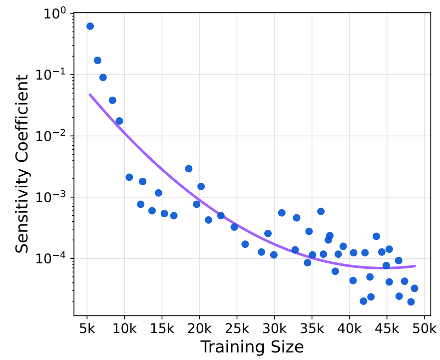

Figure 5: Empirical estimate of $\delta_{\hat{T}_\phi}$ versus training size.

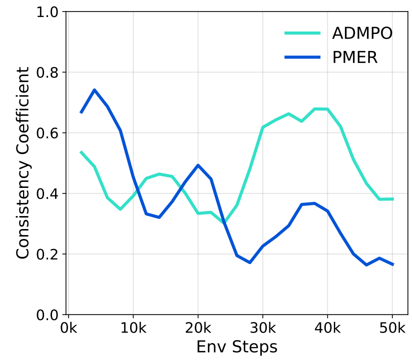

Figure 6: Empirical estimate of $L_s^{(k)}(\hat{T}_\phi)$ on Hopper. We estimate the k=10 consistency coefficient by discretizing the state space into bins, constructing the policy-induced closed-loop transition kernel using the learned dynamics model, computing its k-step kernel P^k, and evaluating $\delta(P^k) = \frac{1}{2}\max_{i,j} \left\| P^k(i,\cdot) - P^k(j,\cdot) \right\|_1$. Lower values indicate better finite-horizon multi-step consistency.

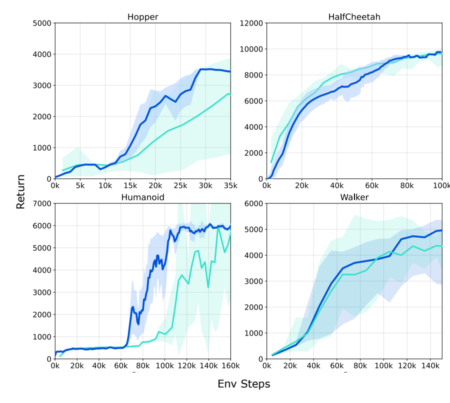

Figure 7: Comparison between PMER and MaPER-MBPO on the four benchmark environments used in MaPER.

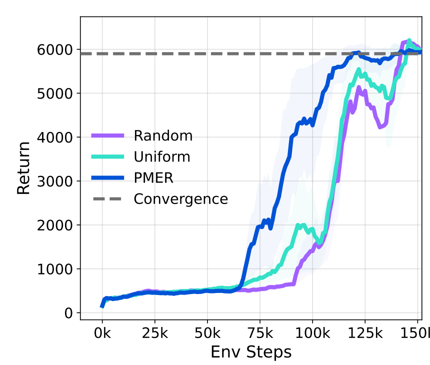

Figure 8: Comparison of PMER, the uniform baseline, and a random-weight proxy on Humanoid task.

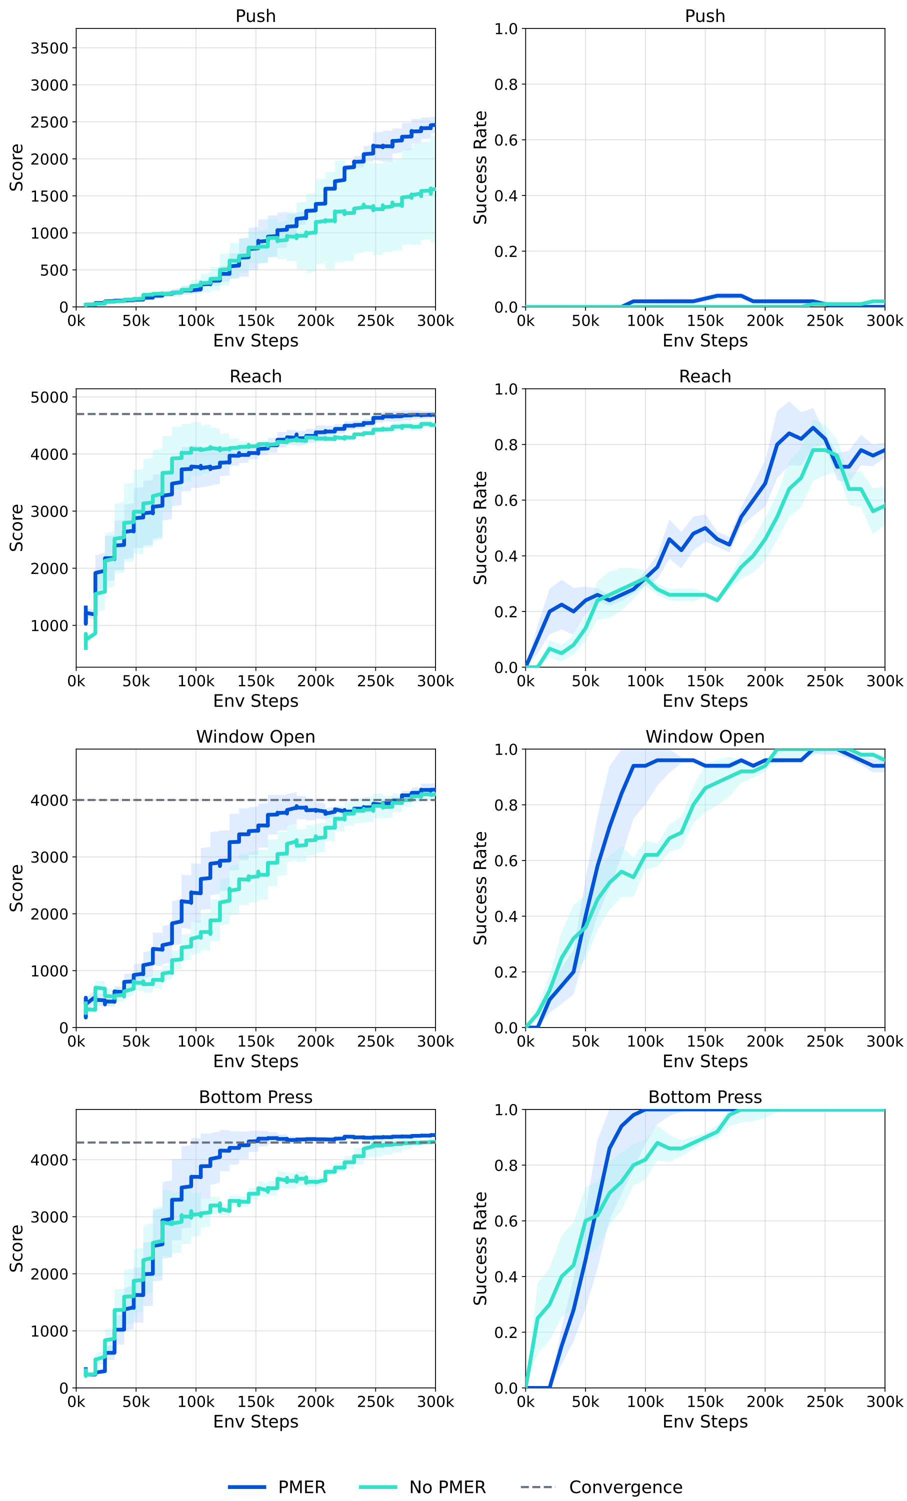

Figure 9: Meta-World results with a Dreamer-v3 backbone.